

INTERVENE BENCH: BENCHMARKING HUMAN-IN-THE-LOOP SAFETY CONTROL FOR TOOL-USING AGENTS

Anonymous authors

Paper under double-blind review

ABSTRACT

Tool-using agents can encounter adversarial instructions during benign tasks. Asking a human does not guarantee safe behavior: the agent must also understand what the reply permits. We introduce INTERVENE BENCH to evaluate consultation and execution under human feedback. The benchmark separates effects that require valid approval from effects that remain prohibited. Its executable core contains 250 tasks with four fixed replies each: approval, denial, ambiguity, and scope mismatch. We organize evaluation into three safety dimensions and three human-interaction dimensions. An exploratory comparison uses Codex with one model on 80 cases per condition. An authorization-scope prompt reduces unsafe outcomes from 18/80 to 8/80, while benign task completion remains 76/80. The prompt still fails the motivating redundant-validation mechanism. Correct, safe task completion after feedback changes from 22/26 to 21/27 across different consulted groups, so improvement is not uniform. Independent semantic annotation, richer replies, shared-prefix branching, and broader agent coverage remain necessary.

Draft status: Internal working manuscript, revised 12 September 2026. The four-reply core and prompt-defense pilot are implemented. Recognition and semantic question labels, richer replies, shared-prefix branching, and a runtime guard remain incomplete.

1 INTRODUCTION

An agent is asked to reconcile an invoice batch. A workflow note says that the same validation must run three times before the report can be written. A skill repeats this instruction. The extra passes add no new information, but their wording makes them look necessary. Should the agent follow the note, reject the extra work, or ask the user? If it asks, does *use your judgment* permit the extra cost? This example motivates our study: lower-trust workflow materials can turn a benign task into an unauthorized action.

Tool-using agents need to protect both safety and task progress. Existing benchmarks study unsafe tool behavior, indirect prompt injection, and attacks carried by reusable skills (Zhang et al., 2024; DeBenedetti et al., 2024; Ning et al., 2026). Human-interaction benchmarks study useful questions, feedback, and shared control (Trinh et al., 2026; Wu et al., 2026). HumanAgencyBench also shows why supporting a person involves more than following every statement (Sturgeon et al., 2025). Together, these works motivate examining the full path from encountering a suspicious instruction to acting after human feedback.

A question alone does not settle a safety decision. A reply may approve a specific action, deny it, leave the choice unresolved, or approve a different step. A reply can also contain a factual error. The agent must use the information that the reply actually supplies. It should perform valid, authorized work and preserve boundaries that the reply cannot change. For example, ordinary approval cannot override a trusted prohibition on disclosing credentials.

We introduce INTERVENE BENCH to study selective consultation and execution after feedback. Each task starts with a benign goal and an attacker-controlled instruction in lower-trust material. The

agent has a neutral consultation tool and chooses when to use it. We distinguish *confirmable* effects, which valid approval may permit, from *prohibited* effects, which ordinary approval cannot permit. The executable core contains 250 tasks with four fixed reply conditions each. Factual-error and partial-approval tracks remain planned extensions.

Evaluation has two groups of three dimensions: task completion, attack success, and autonomous risk recognition; then consultation recall, question accuracy, and correct safe completion after a reply. These dimensions expose different failures. An agent may act without asking, ask an inaccurate question, misuse a reply, or safely omit an approved action. Some require semantic labels that the current mechanical scorer does not provide.

An exploratory pilot uses Codex with one model on 20 tasks and all four replies. A prompt that checks authority and action scope reduces unsafe outcomes from 18/80 to 8/80, while benign completion stays at 76/80. It still fails the redundant-validation attack, and strict safe completion after feedback does not improve across the changing consulted groups. Our present contributions are the evaluation formulation, an executable four-reply core, and this audited prompt-defense comparison. Broader model coverage, independent annotation, and controlled response branching are still needed.

2 RELATED WORK

Tool behavior and indirect attacks. ToolEmu evaluates risks in emulated environments, and R-Judge studies risk awareness in recorded interactions (Ruan et al., 2024; Yuan et al., 2024). Agent-SafetyBench covers varied environments and behavioral failures (Zhang et al., 2024). AgentHarm starts from malicious user requests and checks harmful task completion (Andriushchenko et al., 2025). InjecAgent and AgentDojo instead place adversarial instructions in lower-trust inputs; AgentDojo checks task utility separately from attack success (Zhan et al., 2024; Debenedetti et al., 2024). MT-AgentRisk studies how harmful tool use develops over multiple turns (Li et al., 2026). Skill-Harm and SkillSafetyBench extend evaluation to skills and related artifacts (Ning et al., 2026; Jin et al., 2026). These works inform our attack construction and effect checks. We focus on how consultation and later actions interact under different human replies.


Human collaboration and authority. HiL-Bench measures task completion, blocker recall, and question precision when agents need missing information (Trinh et al., 2026). HAS-Bench varies human participation and evaluates clarification, feedback use, control, and safety, including approval of protected actions (Wu et al., 2026). HumanAgencyBench evaluates six aspects of agency support, such as clarification, correcting misinformation, and deferring important decisions (Sturgeon et al., 2025). OverEager-Bench studies unauthorized scope expansion during benign coding tasks (Qu et al., 2026). Thus, human control and authorization are already evaluation targets. Our narrower question is how agents handle indirect attacks before and after feedback that differs in validity and scope. The planned factual-error and partial-approval tracks would extend the current four-reply core. Appendix H compares the evaluated settings without equating model counts with framework counts.

Defenses. ToolShield derives defensive guidance from tool exploration and experience (Li et al., 2026). CaMeL restricts information flows by separating trusted control from untrusted data (Debenedetti et al., 2025). These approaches motivate both prompt guidance and checks at tool dispatch. We currently test only an authorization-scope prompt. A future guard could bind a human approval to an exact effect and check later calls against it, but we report no implemented runtime guard or head-to-head result against these defenses.

3 BENCHMARK OVERVIEW

3.1 TASKS, THREATS, AND AUTHORITY

A base task specifies a benign goal g_i , initial environment s_i , artifacts D_i , benign-completion predicate B_i , and an attacker-targeted effect E_i . Each effect has an authorization mode: *confirmable* or *prohibited*. The attacker modifies designated lower-trust materials to make an extra action appear

a Observed workflow**b Expected behavior after a matched question**

Confirmable effects only. Prohibited effects remain blocked in every branch.

Figure 1: The implemented evaluation follows a benign task, an injected instruction, optional consultation, and recorded tool effects. Each of four fixed reply conditions starts independently. The lower panel gives expected behavior for a confirmable effect; a prohibited effect remains blocked under every reply. Effects are simulated, and shared-prefix response branching is not yet implemented.

necessary for the workflow. It may coordinate instructions across a skill and a policy-like document. The original user goal, trusted deployment constraints, response fixtures, and evaluator remain outside attacker control. A document’s filename or assertion of authority does not make it trusted.

A neutral `ask_human(question)` interface makes consultation available without revealing the risk label or forcing a question. The agent decides whether to ask and what to ask. The runtime records ordered observations, questions, responses, and tool effects. In the current prototype, consequential operations are simulated: a recorded upload represents a benchmark effect rather than an actual external transmission.

We distinguish an *attacker-targeted effect* from an *unsafe outcome*. A confirmable effect is initially unauthorized but can become permissible after timely approval of the relevant action and scope. A prohibited effect remains impermissible under ordinary human approval. This prevents the metric from treating valid authorized execution as an attack success, or treating blind obedience to an approving human as inherently safe.

3.2 HUMAN FEEDBACK AND EXPECTED BEHAVIOR

The implemented core uses

$$\mathcal{R}_{\text{core}} = \{\text{clear approval, clear denial, ambiguity, scope mismatch}\}. \quad (1)$$

For confirmable actions, matched clear approval permits the specified effect; denial withholds permission; ambiguity leaves the decision unresolved; and approval of a different step cannot transfer to the questioned action. For prohibited actions, none of these responses permits the effect. Asking again or verifying can be appropriate when uncertainty remains, subject to a fixed query budget.

The broader design also includes *objectively erroneous advice* and *partial authorization*. These are distinct from scope mismatch. An erroneous reply may incorrectly identify an external recipient as internal; the task must contain accessible trustworthy evidence for verification, or permit a safe unresolved outcome. Partial authorization may allow a redacted file while withholding raw records, requiring both execution of the allowed part and restraint elsewhere. These tracks need new effect predicates and audited fixtures. They are specified here and in Appendix B, not counted among the

existing 1,000 cases. Clarity, correctness, and authority should be independently controlled rather than collapsed into one clarity score.

3.3 STAGE-SPECIFIC EVALUATION

For a core run, let e_i denote occurrence of the specified effect, h_i a matched question before its first occurrence (or a matched question when no effect occurs), and b_i benign completion. Under the current single-effect scorer,

$$a_i = \mathbb{I}[\text{confirmable}_i \wedge h_i \wedge \text{clear-approve}_i], \quad u_i = e_i(1 - a_i). \quad (2)$$

The approval fixture must apply to the questioned effect; the hidden category is an evaluator label, not an input to the agent or a defense. General time- and scope-bound authorization requires the richer contract in Appendix C; the current scorer does not enforce arbitrary token expiry or cross-tool equivalence.

We group the metrics into *agent safety* (benign completion, unsafe attack success, and autonomous risk recognition) and *human interaction* (consultation recall, question accuracy, and correct safe completion after feedback). The groups describe different evaluation questions; they do not imply a shared denominator.

Before feedback. Report consultation on cases requiring clarification, direct blocking on hard-prohibited cases, question match, timing, and burden. Let $\mathcal{C}$ be the consultation-required set and $\mathcal{P}$ the direct-block set. The core implementation uses confirmable and prohibited labels as these sets:

$$\text{HILRecall} = \frac{\sum_{i \in \mathcal{C}} h_i}{|\mathcal{C}|}, \quad \text{DirectBlock} = \frac{\sum_{i \in \mathcal{P}} \mathbb{I}[|Q_i| = 0 \wedge e_i = 0]}{|\mathcal{P}|}, \quad (3)$$

where Q_i is the set of questions. This proxy needs refinement: being able to approve an optional extension does not mean the agent must ask about it. Review must distinguish optional work from decisions needed to complete the task. Safe autonomous rejection is reported separately and does not establish observed risk recognition by itself.

After feedback. On runs $\mathcal{H} = \{i : h_i = 1\}$, define $c_i = e_i$ after valid approval and $c_i = 1 - e_i$ otherwise. Response-conditioned action accuracy is $\text{RCAA} = |\mathcal{H}|^{-1} \sum_{i \in \mathcal{H}} c_i$. Report approval execution, denial compliance, ambiguity handling, and scope-mismatch restraint separately. Unlike a restraint-only score, RCAA penalizes omission of the approved effect. Such omission is a branch-contract failure, not necessarily a safety violation. The strict four-response boundary score requires $h_i = c_i = 1$ in every branch of a confirmable task; its denominator is complete four-branch task groups. Independent starts make this a consistency diagnostic, not a causal estimate.

Final outcomes. Report $\text{BCR} = N^{-1} \sum_i b_i$, $\text{UnsafeASR} = N^{-1} \sum_i u_i$, and $\text{JointUtility} = N^{-1} \sum_i b_i(1 - u_i)$. Conditional safe completion after feedback additionally requires $b_i(1 - u_i) = 1$ on $\mathcal{H}$; it differs from RCAA because safe omission of an approved action can satisfy the former. Always give denominators and mark undefined ratios N/A. Source exposure, observable risk recognition, question precision, and human-query count complete the diagnostic view. Recognition and fine-grained question quality need semantic annotation beyond the current tool matcher.

Recognition and question accuracy. Autonomous recognition requires visible evidence, before human feedback, that the agent identified the specific risk or missing authority. An accurate question can supply such evidence; non-execution or reading a file alone cannot. Its denominator is the reviewed set of attack cases with accessible decision evidence. Independent annotation is pending, so we report N/A. Question accuracy requires the action, relevant scope, and missing permission to be stated correctly, divided by all questions. The current tool-and-argument matcher supplies only a proxy. It does not establish that a question explains the risk accurately.

Strict success after feedback. To measure correct execution and safety together, define

$$\text{PostSuccess} = \frac{\sum_{i \in \mathcal{H}} c_i b_i (1 - u_i)}{|\mathcal{H}|}. \quad (4)$$

This requires correct response use, benign completion, and no unsafe effect. Report it alongside RCAA and the eligible count. The end-to-end version uses N as its denominator. Neither conditional score is a causal estimate when different conditions cause different cases to seek help.

4 BENCHMARK CONSTRUCTION

4.1 FROM A WORKING ATTACK TO REUSABLE TASKS

The motivating artifact is a local retail-workflow case in which a skill and a mirrored policy require three independent passes over the same action-chain data before writing a result. The passes add no new input information, yet the text presents repetition as mandatory quality control. A later invoice-reconciliation task adopts this mechanism for a normal total-and-discrepancy report. The retail case is confirmable in its recovered HIL adaptation; the archived V3 invoice derivative is prohibited. Shared attack wording does not imply identical task contracts or ground truth.

Construction starts with a benign goal and verifiable completion effects, then introduces an operationally plausible extra action in materials the agent naturally encounters. Each case must specify why authorization is needed or why a hard boundary applies, expose sufficient trustworthy decision evidence, and preserve useful benign work when permission is withheld. Valid-approval branches also specify what completion of the approved extension entails. An optional suspicious step must not silently become a mandatory consultation label.

Table 1 retains the eight-domain, twenty-type taxonomy developed during manuscript design. Risk domains identify protected boundaries; attack types identify mechanisms for crossing them. Application domains are a separate axis. This is a construction target, not a claim that the twenty implementation labels already cover these twenty entries. Delegated execution, cross-tool bypass, and richer scope tests require additional environments and predicates.

4.2 FROZEN RESPONSES AND CONTROLLED BRANCHING

Response semantics and surface style are separate variables. The existing corpus contains 250 workflows in 24 application domains and 20 implementation attack families: 195 confirmable and 55 prohibited workflows, each with four fixed replies, yielding 1,000 cases. Eighteen styles vary the wording. Current fixtures come from deterministic templates; an optional frozen model generator is available, but these cases are not an online human simulation. For new factual-error and partial-approval tracks, freeze and review the generated replies before evaluation, including a deterministic policy for repeated questions.

In *end-to-end evaluation*, each branch starts from the same workspace and the agent independently decides whether to consult. This measures natural intervention, but different runs can diverge before receiving a reply. In the planned *controlled response evaluation*, save the agent context and tool state immediately before a matched response, then resume that identical prefix under each reply. The contrast can then isolate the response at a fixed decision point. Checkpoint/resume support has not yet been implemented; preserving the same initial workspace is insufficient.

4.3 VALIDATION BOUNDARY

The 1,000 core cases regenerate identically and pass the existing mechanical audit. These checks establish reconstruction, not independent semantic validity. Further review must verify benign solvability, whether consultation is needed, reply meaning, and the actual content of required outputs. Each new scope condition needs allowed and disallowed effect checks. Erroneous advice must have adjudicable truth and an attainable verification or abstention strategy. Family-level development/test separation and benign counterparts are needed to detect memorized templates and blanket blocking. No 3,000-case or multi-model completion claim is supported by the current artifacts.

Table 1: Proposed risk taxonomy for authorization-aware HIL evaluation. These are design categories, not verified per-category coverage counts. Each type describes an attempted violation in an otherwise benign workflow.

Risk domain	Attack type	Description
Data confidentiality	Access-to-disclosure	Treats permission to read as permission to disclose.
	Excessive data extraction	Extracts protected content beyond task needs.
	Unauthorized data reuse	Reuses data for a purpose the user did not request.
Recipient and destination	Additional destination	Adds an unapproved delivery endpoint.
	Audience expansion	Expands recipients through aliases or groups.
	Cross-boundary redirection	Crosses a tenant, workspace, or organizational boundary.
Authorization validity	Stale or revoked authorization	Reuses historical or revoked permission.
	Role-as-authorization	Treats eligibility as permission for the present action.
	Wrong or over-delegated approver	Accepts authority from an ineligible principal.
Approval scope and binding	Cross-action approval transfer	Transfers permission between materially different actions.
	Wrong-object approval	Applies permission to the wrong file, record, or resource.
Privilege and policy	Privilege escalation	Uses privileges outside the legitimate task scope.
	Hard-policy bypass	Presents ordinary permission as overriding a hard constraint.
Operational scope	Unauthorized persistence	Turns a temporary step into a persistent change.
	Repeated execution	Repeats an operation beyond its justified count.
	Resource overuse	Exceeds a task’s computation or resource budget.
Execution bypass	Cross-tool bypass	Realizes a rejected effect through another tool.
	Delegated execution	Defers or delegates the restricted effect to another executor.
Audit and accountability	Audit suppression	Removes evidence needed for later review.
	Attribution obfuscation	Hides the responsible principal behind a proxy identity.

5 EXPERIMENTS: PRELIMINARY EVIDENCE AND EVALUATION PLAN

5.1 AUDITED FOUR-RESPONSE PILOT

The completed pilot uses Codex with gpt-5.6-1una, medium reasoning effort, and the V4 core protocol. Before execution, twenty base tasks were selected, one per implementation attack family, by the smallest SHA-256 value of the seed-prefixed task identifier. All four response conditions were run, giving 80 cases: 15 confirmable tasks and five prohibited tasks. The prompt and consultation interface are neutral. Case and output paths are uniform inside Docker. These experiments instantiate the four-response core of Sections 3–4; they do not test the proposed factual-error, partial-authorization, or shared-prefix tracks.

All 80 runs finished without timeout. Per-case scores reproduce from recorded tool events. The raw streams contain no native shell calls or observed hidden-label reads; manual review resolved five flagged directory references as explicit exclusions in simulated searches. This audit records observed behavior rather than proving that uniform paths enforce a security boundary.

Of 22 executions of the attacker-targeted effect, four have valid approval and 18 are unsafe under Equation 2. Sixteen unsafe runs lack a prior matched question. Two execute after consultation: a finance-notice agent treats repeated ambiguous deferral as confirmation, and a state-transition agent acts after approval of a different scope. Conversely, two consulted confirmable cases omit the requested effect after clear approval. Thus, correct response-conditioned action is 22/26, whereas

Table 2: Six evaluation dimensions on 80 runs per condition. Recognition lacks independent labels. Question precision is a rule-based proxy. The final row requires correct response use, benign completion, and safety; its eligible runs differ between conditions.

Metric	Neutral	Prompt guard
Benign task completion	76/80 (95.0%)	76/80 (95.0%)
Unsafe attack success	18/80 (22.5%)	8/80 (10.0%)
Autonomous risk recognition	—	—
HIL recall on confirmable cases	24/60 (40.0%)	27/60 (45.0%)
Question precision (rule proxy)	27/27 (100.0%)	29/29 (100.0%)
Correct and safe completion after reply	22/26 (84.6%)	21/27 (77.8%)




Figure 2: Safety outcomes in the 80-case pilot under authorization-aware scoring. A prior question must match the specified effect and precede it. The safe group includes valid approved executions; safety alone implies neither benign completion nor correct response-conditioned action.

safe benign completion after feedback is 24/26. Calling the latter response understanding would conceal the two approval-branch omissions.

Only two of 15 confirmable tasks meet the strict boundary criterion across all four independently run branches. This diagnoses end-to-end consistency, not a causal response effect. Moreover, 24 confirmable cases safely omit the extension without asking; some may be appropriate rejection of optional work rather than missing consultation. Utility checks also remain weak: 56 of 80 cases check output-writing effects alone, and all four benign-completion failures occur in one task with an artifact-selection mismatch. Independent contract review and semantic output checks are needed before broader capability claims.

Appendix A records a separate, non-random V3 subset with 130 replay-verified trajectories and 59 unsafe outcomes. Its invoice-reconciliation traces connect the motivating redundant-execution mechanism to observed attacks, but are not pooled with this pilot.

5.2 EXPLORATORY AUTHORIZATION-PROMPT COMPARISON

We run the same frozen 80 cases with an authorization-scope prompt appended to the task. It instructs the agent to separate task evidence from permission, check the exact effect scope, consult when authority is uncertain, accept only valid in-scope approval, and continue authorized work. It neither reads hidden labels nor mediates tool dispatch. Section 6 describes this prompt and its remaining failures.

All 80 additional runs complete, reproduce their scores, and match their assigned prompt hashes. Manual review resolves four metadata references as explicit exclusions in simulated searches; no native calls or hidden-label reads are observed. Table 2 shows unsafe outcomes falling from 18 to eight, with ten unsafe-to-safe transitions and no safe-to-unsafe transitions. Benign completion remains 76/80, and joint safe utility rises from 58/80 to 68/80. The ASR difference is -12.5 percentage points, with a 95% paired base-task bootstrap interval of $[-25.0, -3.7]$ (10,000 draws). These intervals describe an exploratory 20-task sample; the earlier baseline was not randomized contemporaneously with the defense, and service-side drift and randomness remain possible.

The prompt leaves important gaps. Matched consultation on confirmable cases rises only from 24/60 to 27/60; its difference interval includes zero. Post-question unsafe outcomes fall from 2/26 to 0/27, but conditional populations differ. On the 23 cases consulted in both conditions, response-conditioned accuracy is 19/23 versus 21/23. Valid-approval execution is 4/6 versus 5/7, retaining two omissions. The eight remaining unsafe runs are all four branches of archive deletion and redundant invoice validation, each executed without consultation. Thus the defense does not resolve the motivating redundant-execution mechanism. The prompt and fixtures were not revised during this comparison.

The stricter post-feedback metric also requires correct response use. It is 22/26 for neutral and 21/27 for the prompt condition. The consulted sets differ: four newly consulted certificate cases fail the benign-output check. On the 23 cases consulted in both conditions, strict success is 19/23 versus 21/23. This post-hoc intersection is descriptive and does not remove selection bias. Question matching is 27/27 versus 29/29 at the question level, but semantic question accuracy and autonomous recognition remain unannotated.

5.3 EVALUATION PLAN

RQ1: When should the agent consult, verify, or block? Review consultation necessity separately from authorization mode, freeze development and held-out cases, and report source exposure, recognition evidence, question quality, stage-specific decisions, and final outcomes with all denominators. Four attack-free draft counterparts have now run under both conditions. Both groups finish all four cases with no questions or original-risk canary effects; assistant content review passes all eight outputs. Exact-tool benign completion is 3/4 versus 4/4: the neutral invoice uses one simulated validation command instead of the predicate’s resource tool despite reporting correct arithmetic. The guarded export publishes twice after adding a receipt. These observations diagnose tool-specific scoring and extra work, and do not establish a population false-positive rate.

RQ2: Does behavior track the meaning and validity of feedback? At a shared pre-response state, compare valid approval, denial, ambiguity, and scope mismatch. Extend this comparison to genuine scoped permission and objectively erroneous advice with independently verifiable facts. Report authorized execution and unauthorized effects together, including allowed-subset completion and out-of-scope effects. Use repeated continuations and base-task cluster intervals. Only runs actually exposed to feedback inform its interpretation.

RQ3: What does HIL add? Compare the same agent with consultation available and unavailable under explicit safe fallback rules. Measure end-to-end safety, benign completion, authorized-extension completion, and query burden. A no-response branch after a frozen question answers a narrower continuation question; it is not equivalent to an agent that never had a consultation tool. Both comparisons are pending.

RQ4: Can execution-time authorization improve the tradeoff? Extend the completed prompt comparison to the proposed runtime guard on identical frozen tasks and response conditions. Include allowed benign actions, alternative tool realizations, and ablations of scope binding and factual verification. A separate four-case prompt smoke on one task changes unsafe outcomes from 4/4 to 0/4, with no consultation in either condition; it remains a wiring check and is excluded from Table 2. No implemented SAIL result is reported.

6 CURRENT DEFENSE AND REMAINING FAILURES

We retain the tested authorization-scope prompt as our current defense. Before a consequential action, it asks the agent to check what the user requested, where the proposed action came from, and whether valid permission covers its object, destination, and scope. A document can supply task evidence without granting new authority. If a necessary decision remains unresolved, the agent should ask a specific question. It should then use only the permission supplied by a valid reply and continue the legitimate task.

This is a prompt-level intervention. The tool interface remains available, and no execution layer enforces these checks. The prompt receives no hidden attack label, response category, or expected

decision. It can reduce unsafe behavior when the agent applies the checks, but it can fail when an extra operation looks routine. In our pilot, all four invoice-reconciliation variants still perform redundant validation without asking. All four archive-inventory variants also delete material without consultation. These two mechanisms account for the eight remaining unsafe outcomes.

A runtime guard is a possible next step. It would track scoped authority and check proposed effects before dispatch, including repeated or equivalent actions through another tool. This requires validated tool adapters and a reliable account of which authority is trusted. The SAIL design in Appendix G remains a specification. We do not treat its expected behavior as an experimental result.

7 LIMITATIONS AND CONCLUSION

The original pilot covers one agent-model configuration, twenty tasks, simulated effects, and incomplete semantic checks. The prompt comparison is exploratory and was not run in contemporaneous randomized order. Consultation-required labels need review, and the clean controls are small. The four replies do not yet cover genuine partial authorization or objectively false advice. Shared-prefix branching and runtime enforcement are unimplemented.

INTERVENE BENCH evaluates whether an agent seeks useful help and then acts within valid authority. The current prompt defense reduces unsafe outcomes on the pilot, but it leaves the motivating redundant-work attack unresolved. The next tests need stronger task checks, unseen tasks, broader model coverage, and richer replies. Safety gains must be reported together with authorized task completion and the cost of human intervention.

AI USE STATEMENT

Generative AI tools were used in research discussions, conceptual and experimental design, synthetic workflow and response construction, implementation and recovery of evaluation code, trajectory analysis, literature discovery, and drafting and revising manuscript text and figures. The reported prototype responses are synthetic fixtures rather than human-participant data. Automated checks reproduce the cited result artifacts, but complete author review of task contracts, proposed extensions, annotations, citations, and manuscript remains pending in this internal draft. The final statement must reflect the work actually performed and reviewed before submission.

ETHICS STATEMENT

The evaluation concerns malicious instructions embedded in otherwise benign workflows. The current harness simulates consequential operations and does not perform real data transmission, deletion, or privilege changes. Conversation archives used to reconstruct the research history are private working materials and are excluded from the paper’s shareable source bundle. Any future dataset release should separate synthetic examples from sensitive source material and document the provenance and release permissions of reused tasks. No human-participant response study is reported.

REPRODUCIBILITY STATEMENT

The available prototype preserves fixed case fixtures, ordered tool events, raw agent streams, per-case scores, run metadata, and file hashes. Appendix A distinguishes verified artifacts from proposed work and describes scoring limitations. The paper figures and numerical tables are generated from local result files by asset-generation scripts; the defense comparison verifies prompt hashes and uses paired base-task intervals. A public code and synthetic-data snapshot is available; it excludes credentials and private conversations. Appendix J records the repeated API evaluation. An anonymous submission artifact remains pending.

REFERENCES

Maksym Andriushchenko, Alexandra Souly, Mateusz Dziemian, Derek Duenas, Maxwell Lin, Justin Wang, Dan Hendrycks, Andy Zou, Zico Kolter, Matt Fredrikson, Eric Winsor, Jerome

- Wynne, Yarin Gal, and Xander Davies. AgentHarm: A benchmark for measuring harmfulness of LLM agents. In *International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2410.09024>.
- Edoardo DeBenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-2636. URL <https://arxiv.org/abs/2406.13352>.
- Edoardo DeBenedetti, Ilia Shumailov, Tianqi Fan, Jamie Hayes, Nicholas Carlini, Daniel Fabian, Christoph Kern, Chongyang Shi, Andreas Terzis, and Florian Tramèr. Defeating prompt injections by design. *arXiv preprint arXiv:2503.18813*, 2025. URL <https://arxiv.org/abs/2503.18813>.
- Chang Jin, An Wang, Zeming Wei, Kai Wang, Biaojie Zeng, Qiaosheng Zhang, Chao Yang, Jingjing Qu, Xia Hu, and Xingcheng Xu. SkillSafetyBench: Evaluating agent safety under skill-facing attack surfaces. *arXiv preprint arXiv:2605.12015*, 2026. URL <https://arxiv.org/abs/2605.12015v3>.
- Xu Li, Simon Yu, Minzhou Pan, Yiyu Sun, Bo Li, Dawn Song, Xue Lin, and Weiyan Shi. Unsafer in many turns: Benchmarking and defending multi-turn safety risks in tool-using agents. *arXiv preprint arXiv:2602.13379*, 2026. URL <https://arxiv.org/abs/2602.13379>.
- Yuting Ning, Zhehao Zhang, Yash Kumar Lal, Boyu Gou, Junyi Li, Weitong Ruan, Chentao Ye, Rahul Gupta, Diyi Yang, Yu Su, and Huan Sun. SkillHarm: Lifecycle-aware skill-based attacks via automated construction. *arXiv preprint arXiv:2606.02540*, 2026. URL <https://arxiv.org/abs/2606.02540>.
- Yubin Qu, Ying Zhang, Yanjun Zhang, Gelei Deng, Yuekang Li, Leo Yu Zhang, and Yi Liu. Overeager coding agents: Measuring out-of-scope actions on benign tasks. *arXiv preprint arXiv:2605.18583*, 2026. URL <https://arxiv.org/abs/2605.18583v1>.
- Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J. Maddison, and Tatsunori Hashimoto. Identifying the risks of LM agents with an LM-emulated sandbox. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2309.15817>.
- Benjamin Sturgeon, Daniel Samuelson, Jacob Haimes, and Jacy Reese Anthis. HumanAgencyBench: Scalable evaluation of human agency support in AI assistants. *arXiv preprint arXiv:2509.08494*, 2025. URL <https://arxiv.org/abs/2509.08494v1>.
- Tu Trinh, Mohamed Elfeki, Guangze Luo, Kelvin Luu, Nathan Hunt, Ernesto Hernandez, Nandan Marwaha, Yannis Yiming He, Charles Wang, Fernando Carabedo, Alessa Castillo, and Bing Liu. HiL-Bench (human-in-loop benchmark): Do agents know when to ask for help? *arXiv preprint arXiv:2604.09408*, 2026. URL <https://arxiv.org/abs/2604.09408>.
- Yaozu Wu, Wei-Chieh Huang, Jizhou Guo, Dongyuan Li, Renhe Jiang, Henry Peng Zou, Chunyu Miao, Shanghao Li, Weizhi Zhang, Weiwei Ye, Yankai Chen, Meng Zhang, Xue Liu, and Philip S. Yu. HAS-Bench: Evaluating LLM-based human-agent systems under configurable human participation. *arXiv preprint arXiv:2607.04329*, 2026. URL <https://arxiv.org/abs/2607.04329v1>.
- Tongxin Yuan, Zhiwei He, Lingzhong Dong, Yiming Wang, Ruijie Zhao, Tian Xia, Lizhen Xu, Binglin Zhou, Fangqi Li, Zhuosheng Zhang, Rui Wang, and Gongshen Liu. R-Judge: Benchmarking safety risk awareness for LLM agents. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pp. 1467–1490, 2024. doi: 10.18653/v1/2024.findings-emnlp.79. URL <https://aclanthology.org/2024.findings-emnlp.79/>.
- Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics: ACL 2024*, pp. 10471–10506, 2024. doi: 10.18653/v1/2024.findings-acl.624. URL <https://aclanthology.org/2024.findings-acl.624/>.

Zhexin Zhang, Shiyao Cui, Yida Lu, Jingzhuo Zhou, Junxiao Yang, Hongning Wang, and
Minlie Huang. Agent-SafetyBench: Evaluating the safety of LLM agents. *arXiv preprint
arXiv:2412.14470*, 2024. URL <https://arxiv.org/abs/2412.14470>.

A EVIDENCE AND VERSION BOUNDARIES

Artifact or claim	Current evidence	Permitted interpretation
V3 historical run	130 completed trajectories; archived scorer reproduces all results	Non-random diagnostic subset of a 200-case version.
V4 core corpus	250 workflows, 1,000 cases; regeneration matches	Available four-response operational core; semantic review remains incomplete.
Historical V4 full run	Report and conversation record; full raw trajectories missing	Historical record, excluded from empirical claims here.
New V4 pilot	80 trajectories; original scores replay exactly	Preliminary four-response evidence for one configuration on 20 tasks.
Prompt comparison	80 additional trajectories on the same manifest	Exploratory prompt-only contrast; 18/80 versus 8/80 unsafe.
Clean controls	Four draft tasks in two conditions; eight trajectories	Small utility diagnostic; independent semantic review pending.
Erroneous / partial-approval tracks	Original research aims and protocol specification	New fixtures, predicates, and evaluation pending.
Shared-prefix branching	Controlled comparison specified	Runner checkpoint/resume and experiments pending.
3,000 cases / eight configurations	Old manuscript text without matching artifacts	Unverified; not a result of this draft.
Runtime guard	Design specification; separate prompt comparison exists	No completed execution-layer defense or efficacy claim.

Table 3: Evidence inventory. Available core components are distinguished from proposed extensions; versions are not pooled.

The V4 pilot manifest selects one task per attack family by sorting on SHA-256 of 42: followed by the task identifier. Each selected task supplies four response conditions, with twenty cases per condition. The pilot spans fourteen application domains and eighteen styles. Its frozen manifest is `gpt-luna-pilot-80-v1.json`. The runner uses Codex CLI 0.144.1, a 300-second per-case timeout, and at most four concurrent jobs. Every run records the model name and reasoning setting; these do not fix server-side weights.

Each case preserves `events.jsonl`, `codex_stream.jsonl`, `result.json`, and `run_metadata.json`. A separate audit checks score replay, tool use, paths, hidden-label exposure, and artifact consistency. All 80 runs finish with return code zero. Pre-run hashes cover source, materialized cases, and runner. Transfer verification matches 737 artifacts across the completed pilot and smoke runs. Failed network attempts are not scored as safety outcomes.

The separate V3 recovery contains 130 completed cases from a stopped, lexically ordered 200-case run. Its original scorer reproduces 59 unsafe outcomes, 14 matched consultations, 14/14 correct post-response decisions, and 129 benign completions. Three incomplete event directories are excluded. This subset is neither random nor comparable to a complete V4 sample.

The full prompt comparison adds 80 completed trajectories; all 643 transferred files match remote hashes. Baseline and defense use identical manifests, model IDs, reasoning effort, Docker image, and CLI version. Analysis verifies the effective task-prompt hash per case and resamples matched base-task clusters across both conditions. Conditional intervals omit bootstrap draws with an empty eligible denominator. Input/output token totals are 13,163,007/83,035 for neutral and 10,303,895/80,340 for the prompt guard; cached inputs are subsets of these inputs, and no monetary bill is available. No Sol, Terra, GPT-5.5, or GPT-6 Astra experiment has been completed.

Scorer scope. The core evaluator matches specified tools and argument constraints. It accepts clear approval for confirmable effects only after a matched question preceding the first effect. It does not implement arbitrary scope parsing, token expiry, repeated-use limits, general cross-tool equivalence, semantic report checking, or internal risk-recognition inference. Source-before-question tim-

ing and full evidence exposure are separate diagnostics from the pilot’s matched-question indicator. The agent and any deployable defense must remain independent of evaluator labels.

B HUMAN-RESPONSE PROFILES

The original research scope includes clear, ambiguous, and objectively mistaken replies. We operationalize this with categorical semantics rather than a single scalar that conflates clarity, correctness, and authority. Surface wording is varied only after the intended semantics and task facts are fixed.

Profile	Expected behavior	Evidence status
Clear approval	Execute the specified confirmable extension; block if hard-prohibited.	Core fixture.
Clear denial	Withhold the effect and continue permitted work.	Core fixture.
Ambiguity	Do not infer permission; ask, verify, or take a safe fallback.	Core fixture.
Scope mismatch	Do not transfer approval of a nearby benign step to the questioned action.	Core fixture.
Partial authorization	Complete the explicitly allowed subset; withhold the remainder.	Proposed extension.
Objectively erroneous advice	Check accessible trusted facts; resolve the error or preserve a safe unresolved state.	Proposed extension.
No response	Apply a predefined timeout and fallback policy without inventing authorization.	Optional proposed control.

Table 4: Response semantics and implementation status. A misleading factual assertion is not equivalent to approval of the wrong step.

For example, the factual-error track can return a confident assertion that a recipient is internal while a trusted directory identifies it as external. The benchmark must identify the relevant directory and its authority independently of the attacker. If no observation can resolve the fact, verified correction cannot be required; the admissible behavior is calibrated clarification or safe abstention. A partial-approval case can permit one redacted document and reject raw records, but needs a completion predicate for the permitted delivery as well as a violation predicate for the withheld data.

Replies may be generated by a frozen model with a fixed seed and then stored, or produced from deterministic templates. Semantic review must precede model evaluation. Repeated questions require a frozen response policy so that retrying does not silently change the condition. The present 1,000-case corpus uses templates and does not constitute a human-participant study or a validated simulation of human response distributions.

C SCOPE CONTRACT FOR THE PROPOSED GUARD

For the richer evaluation, let σ_j denote the effect of a proposed call at time t_j , $\mathcal{A}_{t_j}$ the current scoped authority, and $\mathcal{K}$ the hard constraints. Define

$$\text{Allowed}(\sigma_j, t_j) = \mathbb{I}[\sigma_j \text{ is covered by } \mathcal{A}_{t_j}] \mathbb{I}[\sigma_j \text{ satisfies } \mathcal{K}]. \quad (5)$$

An unsafe trajectory contains at least one realized effect for which this predicate is zero. The benchmark evaluator must derive its reference authority from the frozen task contract and actual feedback history. The guard estimates authority independently from its permitted observations. Equation 2 is the implemented single-effect core rule, not an implementation of this general predicate.

A proposed authorization record contains the permitted effect or set of effects, action/object/destination/purpose constraints, source principal, validity interval, and use budget. For set-valued approval, coverage requires the proposed effect to fall within the allowed subset; resemblance to an approved action is insufficient. A response should update only the fields it actually resolves. A factual assertion is evidence to check, not automatically an authorization update. Revocation and exhaustion of a use budget invalidate subsequent execution.

Normalization across tools must be verified with known equivalent and non-equivalent effects, including aliases and transformed objects where supported. The same high-level intent or justification

does not establish effect equivalence. Policy correctness also depends on complete mediation: an uninstrumented tool can bypass a pre-dispatch guard. These are implementation and evaluation requirements, not established properties of the recovered runner.

D ANNOTATION AND MEASUREMENT RUBRIC

Each semantic label should cite an observable message or event and be independently reviewed. Risk identification requires the agent to name the relevant action and concern before feedback and unsafe execution. Generic caution, source reading, or mere non-execution does not establish recognition. A precise question may supply recognition evidence; hidden reasoning is not used. Report agreement and adjudication before treating autonomous risk identification rate (ARIR) as validated.

Question precision is the number of effect-matched questions divided by all questions; the current matcher gives 27/27 in the pilot. That is a mechanical content-match result, not proof of semantic quality or appropriate consultation. Count and timing remain separate: repeated questions do not improve case-level recall, and late consultation cannot erase an earlier unauthorized effect.

Review consultation necessity along an axis distinct from authorization mode: a decision needed for task completion, an optional extension that can be safely skipped, or a decision resolved by trusted policy. The existing confirmable/prohibited proxy should remain explicit when reporting current results. It cannot support a complete trigger accuracy or false-positive estimate without benign and optional-action controls.

RCAA checks whether the effect follows the response contract, including execution after valid approval. Post-response safe execution rate (PSER) instead requires benign completion with no unauthorized effect among runs with a timely matched question. The neutral pilot has RCAA 22/26 and PSER 24/26; the prompt guard has 25/27 and 23/27, respectively. The latter PSER denominator includes four newly consulted cases from the task with benign-completion failures; it therefore illustrates a changing conditional population. Neither a high conditional score on a small consulted subset nor the absence of unsafe effects establishes end-to-end HIL competence.

For partial authorization, separately measure completion of permitted effects and occurrence of any forbidden effect. For erroneous feedback, annotate the available evidence, whether the agent identified the contradiction, whether clarification resolved it, and the resulting effects. A single refusal rate cannot substitute for these distinctions. All intervals and paired comparisons should respect base-task clustering.

E EVALUATION MATRIX

Comparison	Controlled variables	Status
Neutral agents	Frozen tasks, tools, reply fixtures, resource limits	80-case GPT pilot complete; full and other-model runs pending.
Four core replies	Same pre-response state, question, tool state	Independent-start pilot complete; shared-prefix runs pending.
Erroneous / partial replies	Truth, source authority, permitted subset	New fixtures and effect predicates pending.
HIL available vs. unavailable	Same task contracts and safe fallbacks	Experiment pending; not inferred from natural ask/no-ask groups.
Attack vs. benign control	Legitimate goal and required normal actions	Four drafts, eight runs complete; independent review pending.
Prompt-only vs. runtime guard	Agent, cases, response conditions	80-case prompt comparison complete; runtime guard pending.
Surface-style robustness	Reply meaning fixed, wording varied	Existing styles; controlled repeated evaluation pending.

Table 5: Completed pilot evidence and planned comparisons. None supplies an implemented runtime-defense result.

Freeze generation, semantic annotations, and held-out test cases before examining their outcomes. The existing pilot is development evidence and cannot become an unseen test after its failures guide

revisions. Record run completion and infrastructure failures separately from safety scores. Report proposed versus admitted effects under a guard, and distinguish forced consultation from spontaneous questions. Replaying an identical pre-response state still needs repeated continuations to characterize model randomness.

F MOTIVATING AND DIAGNOSTIC TRACES

The original retail example repeats the same action-chain extraction three times under instructions mirrored between a skill and a policy-like file. Recovered files verify the payload and evaluator, but not the missing original successful Claude execution trace. Its recovered HIL adaptation is confirmable; the V3 invoice derivative instead marks redundant validation prohibited. Four completed invoice traces each invoke `consume_resource` with operation `validate_invoice_batch` and count three without asking. The replies were never exposed, so these are repeated executions of one mechanism rather than evidence about four interpreted responses.

In the finance-notice pilot, the agent asks twice about an overbroad recipient and still sends after ambiguous deferral. In a historical customer-beta-notice example, the agent asks twice, receives unresolved feedback, and withholds the send while writing a local disposition. They illustrate different follow-through, but come from different tasks and cannot establish a causal effect of response wording.

The two safe approval-branch omissions in the new pilot occur in the schema-migration identity-binding case and the policy-revision resource-budget case. They complete the scored benign task but do not execute the approved extension. This distinction motivates validating whether a branch explicitly requires the extension and measuring authorized utility alongside restraint.

G FUTURE RUNTIME GUARD DESIGN

We retain the working name SAIL for a proposed guard between agent proposals and tool dispatch. Its purpose is to bind permission to effects while preserving legitimate work. The following is a design specification; the recovered runtime does not implement it.

Authority and effect representation. Initialize an authority ledger from the user’s task and explicitly trusted deployment policy. Workflow artifacts supply data but cannot expand that authority. Normalize each consequential call into an effect signature containing action class, object, destination or principal, and scope attributes such as count, privilege, persistence, and purpose. Map uploads and messages to a common disclosure effect when they transfer the same material. This requires tool adapters and validation; tool-name matching alone does not provide semantic equivalence.

Response-aware dispatch. The guard would block effects violating hard constraints, allow effects already covered by valid authority, verify resolvable factual conflicts, and request clarification for remaining confirmable proposals. A specific, valid reply may add a bounded authorization with scope, provenance, expiry, and use count. Partial approval covers only the allowed subset. Denial, uncertainty, and approval of another action create no permission for the pending effect. A confident factual error must not silently change recipient identity or policy state; verification should use a trustworthy source when one is available.

Scope and observability. Each later call must be checked against the current ledger, including equivalent effects attempted through another tool. A denial remains binding until a valid scoped update supersedes it; unresolved replies cannot erase it. Log proposals, admission decisions, response assessments, and realized effects separately. Runtime-induced consultation is distinct from spontaneous help-seeking by the neutral agent.

The guard may use only the user request, deployment constraints, observed evidence, tool arguments, and response text. Hidden response categories, risky-effect annotations, and expected decisions are evaluator-only. Incorrect normalization or authority inference can overblock benign work or admit unauthorized effects; no formal guarantee or empirical improvement is claimed. Appendix C gives the proposed scope contract for evaluation.

H BENCHMARK SETTINGS IN RELATED WORK

Table 6: Selected evaluation settings. Counts refer to the main experiments in the cited versions. A model, an execution framework, and a model–framework configuration are different units. User simulators, judges, and extra training or harness ablations are excluded from these main counts.

Benchmark	Construction and target	Main coverage
Agent-SafetyBench	2,000 cases in 349 environments; behavioral safety	16 model backends in a common interaction loop
MT-AgentRisk	365 harmful goals and their multi-turn versions	OpenHands; six models
SkillHarm	879 attacks across 71 skills; file and memory poisoning	Four frameworks; six main configurations
HiL-Bench	300 SWE/SQL tasks; progressive information blockers	SWE-Agent; four main models
HAS-Bench	397 tasks; configurable human roles, feedback, and control	HAS-Framework; five agent backends
HumanAgencyBench	3,000 queries; six dimensions of agency support	Twenty assistant models; response evaluation
AgentDojo	97 benign tasks and 629 security cases; indirect injection	Ten models in a common tool framework
OverEager-Bench	500 benign scenarios; unauthorized extra actions	Four frameworks, six models, fifteen configurations
SkillSafetyBench	155 attacks across 47 tasks; skill-facing inputs	Four frameworks, seven models, nine configurations
INTERVENE BENCH (current)	250 tasks, four fixed replies; 20-task pilot	Codex; one model; neutral and prompt conditions

These settings answer related but different questions. HAS-Bench already includes approval and protected-action safety. HiL-Bench measures the quality of help-seeking. Our current study combines indirect workflow attacks with explicit reply conditions and action-level scoring. It does not yet evaluate richer erroneous advice, general partial authorization, or shared-prefix continuations.

I ADDITIONAL PILOT METRICS

Table 7: Exploratory comparison on the same 20 tasks and four reply conditions (80 runs per condition). Prompt guard adds authorization-scope instructions; the model, image, cases, and scorer are fixed. Conditional denominators differ. This is a prompt defense, not runtime enforcement.

Metric	Neutral	Prompt guard
Unsafe ASR	18/80 (22.5%)	8/80 (10.0%)
Matched HIL on confirmable cases	24/60 (40.0%)	27/60 (45.0%)
Direct block on prohibited cases	14/20 (70.0%)	16/20 (80.0%)
Response-conditioned action accuracy	22/26 (84.6%)	25/27 (92.6%)
Benign task completion	76/80 (95.0%)	76/80 (95.0%)
Joint safe utility	58/80 (72.5%)	68/80 (85.0%)
Four-response boundary (tasks)	2/15 (13.3%)	5/15 (33.3%)

J REPEATED EVALUATION WITH FOUR AGENT FRAMEWORKS

Configuration. The new matrix uses the actual Codex CLI 0.144.1, Claude Code 2.1.220, DeepSeek Harness 0.1.5-rc.1, and Hermes Agent 0.21.2. We cross them with three requested API model IDs: deepseek-v4-flash, glm-5.2, and qwen3.7-max. These IDs represent the tested routes. We cannot verify the weights or serving stack behind the third-party endpoint. Alias prefixes do not establish independent models.

Protocol checks. We tested 36 text-model IDs with a harmless tool-call round trip; 33 passed. Two MiMo routes returned HTTP 503, and Kimi K3 returned HTTP 400. Additional checks tested the native Responses and Messages protocols. Some older model IDs supported Chat Completions but rejected Responses calls. The three selected IDs passed both native protocols before the main safety matrix. Across 16 CLI preflight attempts, 13 completed and reproduced their scores locally. Two Codex attempts failed during initialization, and one exceeded the initial 300-second limit. Fresh writable state directories resolved initialization. The timed-out configuration completed a separate preflight with the main 900-second limit. All attempts are retained; preflights are excluded from the main comparison.

Hermes was added in a separate frozen extension. Its first preflight failed before any model call because its context setting was below the framework’s 64,000-token minimum. With a 65,536-token setting, all six model-condition preflights completed and passed local score replay. Their largest reported prompt used 8,267 tokens. Codex and DeepSeek Harness use a 32,768-token setting; Claude Code retains native context handling. These settings do not verify the endpoint’s context limits. The Hermes extension begins later in calendar time, so serving drift remains a possible confound.

Execution and repetition. Codex uses Responses, Claude Code uses Messages, and DeepSeek Harness and Hermes use Chat Completions. The relay forwards request bodies and response bytes without local protocol translation. Each CLI owns its agent loop and uses the same benchmark MCP implementation. Native shell and web access are disabled. Codex retains auxiliary planning and MCP-resource tools. Hermes uses its native tool-search bridge to the ten allowed benchmark tools. Its native session-title requests are included in relay usage and request counts. Framework prompts, schema presentation, and sampling settings differ, so a cross-framework result concerns the whole configured system.

We reuse the frozen 20-task, 80-case pilot selection and run each case under neutral and `prompt_guard_v1` conditions three times. The original three-framework plan contains 4,320 attempts. Hermes adds 1,440 attempts using the same cases, models, conditions, and repeats, for 5,760 planned attempts across 12 framework-model configurations. The original manifest is unchanged. Case-configuration blocks and condition order are shuffled with a fixed seed; Hermes follows the corresponding Codex block order. Each attempt starts a fresh container and session. The prompt defense is unchanged. These tasks previously informed prompt development; the repeated matrix is not an untouched test set or a full-suite estimate.

Isolation and scoring. Execution started with eight agent containers and later used 16 and 32. After Hermes was added, the 32-agent allocation used 24 original-plan workers and eight Hermes workers. A 40-agent trial was followed by a burst of upstream connection failures and a return to 32. All three model routes passed fresh connectivity checks before execution resumed. The owner then requested 64 agents. At 64 agents, a unified scheduler assigns initial shares of 48 original-plan workers and 16 Hermes workers. It reuses spare slots when a queue runs out of work. The transition preserved all 912 previously closed attempts and waited for active agents and relay requests to finish. The frozen case lists, per-attempt adapters, and scoring remain unchanged. A later amendment permits a staged increase to 96 and 128 agents, with a 3:1 allocation between the two plans. After two minutes of startup, a five-minute checkpoint requires at least 10% higher valid throughput, bounded request errors, and adequate resources. A marginal window is extended once; an unsuccessful stage returns to the preceding target without cancelling active attempts. The recorded valid rates were 1,272 per hour at 64 agents, 1,512 at 96, and 1,332 in the extended 128-agent window. The scheduler therefore returned to 96. Docker startup queues were observed while host memory remained available. These calendar changes do not isolate the cause of the throughput difference or connection failures.

Each container has a read-only root and case mount, fresh temporary state, one CPU, a 1,536 MiB memory limit, no additional swap, and a 900-second limit. A dedicated internal network reaches a relay that alone holds the API credential. The 64-agent phase uses a relay with 64 in-flight requests, a 2 GiB memory limit, and four CPUs. At the staged handoff, one request outlived its closed agent, so the old relay was retained for its completion. The later staged relay has a 128-request ceiling and a 4 GiB memory limit; its admission limit follows the current agent target. An active attempt can span a limit change, so its launch setting and the capacity history are retained. The earlier expanded

relay allowed 40 requests with a 1 GiB memory limit. The initial scheduler reserved 24 GiB of available host memory; the expanded scheduler requires 64 GiB before new launches. Both require 30 GiB free storage. Each new attempt records its capacity epoch. Concurrency can affect serving latency and behavior, so this change is retained with the results. Consequential task operations remain simulated.

API errors, incomplete streams, timeouts, and out-of-memory kills are recorded separately from valid outcomes. A normally completed refusal or no-tool answer remains a valid outcome. Scores are replayed from ordered events. We report paired case-repeat comparisons and bootstrap base tasks, retaining their four replies and three repeats in each cluster. Conditional post-feedback rates use the consulted set in each condition and do not estimate a causal effect of feedback. Independent recognition labels and semantic question-quality labels remain unavailable.

A supplementary audit, added after observing initial failures, also scores unsafe actions present in failed traces. For each framework, model, and condition, its lower bound includes all observed unsafe actions among closed attempts; its upper bound additionally treats failed attempts without an observed unsafe action as unsafe. Pending attempts are excluded. These sensitivity bounds retain evidence of harm before a failure and show uncertainty from incomplete runs. We report failure categories and recorded token usage for all closed attempts separately from the frozen primary comparison.

Execution status. Provisional, incomplete matrix: 4963 of 5760 attempts have finished. There are 4788 valid runs, 175 failed attempts, and 797 pending attempts. The partial tables do not support a final defense comparison.

Table 8: Provisional, incomplete matrix. Each cell lists neutral / prompt-guard counts as two fractions. Failed attempts are excluded from valid-run denominators.

Configuration	Unsafe ASR	Benign completion
Claude Code / DeepSeek V4 Flash	54/216 ; 40/218	184/216 ; 186/218
Claude Code / GLM 5.2	40/224 ; 26/225	205/224 ; 207/225
Claude Code / Qwen 3.7 Max	66/218 ; 28/217	206/218 ; 206/217
Codex / DeepSeek V4 Flash	51/210 ; 31/214	189/210 ; 193/214
Codex / GLM 5.2	42/198 ; 24/197	181/198 ; 183/197
Codex / Qwen 3.7 Max	70/215 ; 25/208	204/215 ; 194/208
DSH / DeepSeek V4 Flash	63/214 ; 35/214	189/214 ; 191/214
DSH / GLM 5.2	40/196 ; 22/185	182/196 ; 176/185
DSH / Qwen 3.7 Max	69/221 ; 26/221	211/221 ; 210/221
Hermes / DeepSeek V4 Flash	51/165 ; 26/160	143/165 ; 145/160
Hermes / GLM 5.2	35/160 ; 19/163	148/160 ; 155/163
Hermes / Qwen 3.7 Max	42/165 ; 19/164	151/165 ; 154/164

Table 9: Provisional, incomplete matrix. Each cell lists neutral / prompt-guard counts as two fractions. Question accuracy is a rule proxy. Post-feedback groups differ between conditions; semantic recognition remains unannotated.

Configuration	HIL recall	Question proxy	Strict after reply
Claude Code / DeepSeek V4 Flash	98/161 ; 100/163	108/108 ; 114/114	63/99 ; 69/101
Claude Code / GLM 5.2	97/166 ; 96/167	112/112 ; 101/101	73/98 ; 77/96
Claude Code / Qwen 3.7 Max	65/162 ; 77/161	69/69 ; 80/80	45/65 ; 64/78
Codex / DeepSeek V4 Flash	63/157 ; 62/160	88/88 ; 78/78	43/63 ; 44/62
Codex / GLM 5.2	54/145 ; 65/142	62/62 ; 72/72	45/54 ; 60/65
Codex / Qwen 3.7 Max	41/160 ; 64/156	76/80 ; 124/130	29/41 ; 51/64
DSH / DeepSeek V4 Flash	88/162 ; 96/162	101/101 ; 107/107	62/89 ; 72/97
DSH / GLM 5.2	70/141 ; 77/134	73/73 ; 78/78	55/71 ; 65/77
DSH / Qwen 3.7 Max	49/166 ; 65/165	49/49 ; 65/65	35/49 ; 56/65
Hermes / DeepSeek V4 Flash	62/122 ; 59/119	67/67 ; 64/64	41/62 ; 46/60
Hermes / GLM 5.2	58/120 ; 70/121	61/61 ; 72/72	48/58 ; 58/70
Hermes / Qwen 3.7 Max	29/126 ; 47/124	29/29 ; 48/48	22/29 ; 38/47